

TRANSACTIONS

- logical units of work in a database
- sequence of operations: read, write, commit, abort
- ensure reliability and consistency of database operations

ACID Properties

- atomicity:
 - a transaction is atomic and treated as indivisible unit
 - either all operations succeed or none do
 - if a failure occurs, changes are rolled back to maintain consistency
- consistency:
 - transactions preserve the database's predefined rules and constraints
 - do not violate integrity constraints
 - violating transactions are rolled back to maintain data integrity
- isolation:
 - transactions are isolated from each other
 - intermediate states of transactions are not visible to others until committed
 - this prevents interference and maintains consistent results
- durability:
 - committed transactions' changes are permanent
 - changes survive subsequent failures or system crashes
 - data persistence is ensured by storing changes reliably

Scheduling Transactions

→ ordering of operations from different transactions

SERIAL SCHEDULE

- transactions executed one after another
- no overlap in execution
- maintains integrity and isolation
- ensures data consistency
- no parallelism

NON-SERIAL SCHEDULE

- allows concurrent execution of transactions
- enables parallelism
- requires careful design for correctness

SERIALIZABILITY

- property of a schedule
- behavior equivalent to a serial schedule

CONCURRENCY CONTROL

- manages concurrent access to the database
- ensures correctness, consistency and serializability
- prevents conflicts and maintains data integrity

CONFLICT SERIALIZABILITY

- condition for serializability
- determines equivalence to a serial schedule
- order of conflicting operations preserved
- final result is equivalent to some serial execution

PRECEDENCE GRAPH

- graphical representation of a schedule
- shows ordering of conflicting operations
- nodes represent transactions
- directed edges represent operation precedence
- no cycles = serializable

$R_1(A) \rightarrow W_2(A)$
 $W_1(A) \rightarrow R_2(A)$
 $W_1(A) \rightarrow W_2(A)$

VIEW SERIALIZABILITY

- condition for serializability
- determines equivalence to a serial schedule
- considers the visibility of data items
- final results is equivalent to some serial execution

view equivalence = two schedules are view equivalent if they produce the same results for all transactions

RECOVERABLE SCHEDULES

- ensures that, if a transaction T_1 reads from transaction T_2 , T_2 must commit before T_1
- prevents dirty reads, only committed data is read

dirty read = occurs when a transaction reads data that has been modified but not yet committed

AVOIDING CASCADING ABORTS

- cascading aborts occur when an initial aborted transaction causes other transactions depending on it to also abort
- techniques used to avoid cascading aborts include strict two-phase locking and timestamp ordering

LOCK-BASED CONCURRENCY CONTROL

- manages concurrent access to shared resources using locks
- ensures data consistency, isolation and integrity

Types of locks

SHARED LOCK

- (s-lock) allows multiple transactions to read a resource concurrently; read-only access

EXCLUSIVE LOCK

- (x-lock) grants exclusive access to resource
- write access, preventing other transactions from reading or writing

Locking protocols

STRICT TWO-PHASE LOCKING (2PL)

- acquires locks in two phases:
 - **graining phase**: acquires locks on data items as needed during the transaction's execution
 - **shrinking phase**: releases locks only after the transaction has completed
- ensures conflict serializability
- prevents unserializable schedules and deadlocks
- holds all locks until the end of a transaction
- prevents cascading aborts and ensures consistency

TWO-PHASE LOCKING (2PL)

- acquires locks in two phases:
 - **graining phase**: acquires locks on data items as needed during the transaction's execution
 - **shrinking phase**: releases locks after the last lock request has been made
- ensures conflict serializability
- prevents unserializable schedules and deadlocks
- allows the possibility of cascading aborts
- locks are released during the transaction's execution, but new locks are not acquired

Deadlocks

- circular waiting between transactions
- can result in a permanent blocking of transactions

PREVENTION

- strategies to avoid deadlocks
- techniques include deadlock detection, timeouts, and resource ordering

DETECTION

- techniques to identify deadlocks in a system
- algorithms include resource allocation and wait-for graph

Joins

Simple Nested Loops Join

$$M + P_E * M * N$$

- cost of scanning E cost of scanning S
→ S is scanned $P_E * M$ times
* E - M pages, P_E records / page *
* S - N pages, P_S records / page *

Page-Oriented Nested Loops Join

$$M + M * N$$

- cost of scanning E cost of scanning S
→ S is scanned M times
* E - M pages, P_E records / page *
* S - N pages, P_S records / page *

Block Nested Loops Join

$$SOT + NBOT * SIT$$

$$NOB = \left\lceil \frac{NPOT}{SB} \right\rceil$$

- scan of outer table scan of inner table
number of blocks in outer table number of pages in outer table
size of block

Index Nested Loops Join

$$M + ((M * P_E) * cFCRS)$$

cost of finding corresponding records in S

* E - M pages, P_E records / page *

* S - N pages, P_S records / page *

SORTING

Simple Two-Way Merge Sort

- pass 0: N runs
- number of passes:
 - $\lceil \log_2 N \rceil + 1$
 - ↳ number of pages in the file to be sorted
- total cost:
 - $2 * \text{number of pages} * \text{number of passes}$
 - $2 * N * (\lceil \log_2 N \rceil + 1)$ I/Os

External Merge Sort

- pass 0:
 - $\left\lceil \frac{N}{B} \right\rceil$ runs \rightarrow number of pages in the file
 - ↳ available buffer pages
- number of passes:
 - $\lceil \log_{B-1} \lceil N/B \rceil \rceil + 1$
 - ↳ number of pages in the input file
- total cost:
 - $2 * N * \left(\lceil \log_{B-1} \left\lceil \frac{N}{B} \right\rceil \rceil + 1 \right)$ I/Os \rightarrow available pages in buffer

Sort-Merge Join

- sorting E cost: $O(M \log M)$
 - sorting S cost: $O(N \log N)$
 - cost of merging: $M + N$ I/Os
 - worst-case scenario: $O(M * N)$ I/Os
- * E - M pages
* S - N pages

Hash Join

→ partitioning cost: $2 * (M + N) \text{ I/Os}$

→ probing cost: $M + N \text{ I/Os}$

→ total cost: $3 * (M + N) \text{ I/Os}$

* E - M pages, P_E records / page *

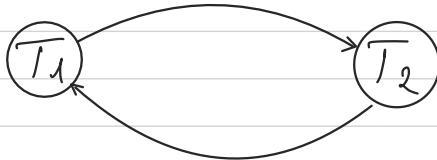
* S - N pages, P_S records / page *

1. Consider the schedule below:

T1	T2
read(A)	
$A = A + 50$	
write(A)	2
	read(A)
	$A = A * 4$
	write(A)
	read(B)
	$B = B * 4$
	write(B)
read(B)	
$B = B + 50$	
write(B)	2



The following integrity constraint must hold: $A = B$. Before the execution of T1 and T2, $A = 100$ and $B = 100$. Is the schedule serializable? Justify your answer.



There is a cycle between T_1 and T_2 which means the schedule is not serializable.

T_1 : read A which is 100

$$T_1: 100 + 50 = 150$$

T_1 : write A with new value 150

T_2 : read A which is 150

$$T_2: 150 * 4 = 600$$

T_2 : write A with new value 600

T_2 : read B which is 100

$$T_2: 100 * 4 = 400$$

T_2 : write B with new value 400

T_1 : read B which is 400

$$T_1: 400 + 50 = 450$$

T_1 : write B with new value 450

The integrity constraint $A = B$ might be violated.

2. Consider the following schedule over transactions T1, T2, T3, T4
(all transactions commit):

•P
(1p)

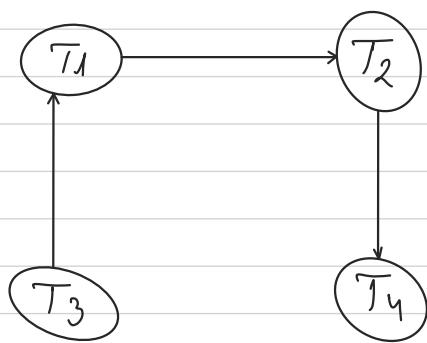
T1	T2	T3	T4
W(A)			R(C)
	R(B)		
		W(D)	
R(D)	R(A)		
R(C)			W(B)

time

$$\begin{aligned} R_1(A) &\rightarrow W_2(A) \\ W_1(A) &\rightarrow R_2(A) \\ W_1(A) &\rightarrow W_2(A) \end{aligned}$$

- a. Draw the precedence graph.
b. Is the schedule conflict serializable? Justify your answer.

a.



b. Because the precedence graph has no cycles, that means the schedule is conflict serializable.
Also, $(W(T1, A), R(T2, A))$ belongs to the conflict relation of S.

4. Encode the data *de gustibus non disputandum* using the secret encryption key *metallica* and the table of codes below. Write the last 5 characters in the result.

	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	u	v	w	x	y	z	-
00	01	02	03	04	05	06	07	08	09	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27

data : *de gustibus non disputandum*

key : *metallica*

m e t a l l i c a
 13 05 20 01 12 12 09 03 01
 d e g u s t i b
 04 05 00 07 21 19 20 09 02
 u s m o m d i
 21 19 00 14 15 14 00 04 09
 s p u t a m d u m
 19 16 21 20 01 14 04 21 13

04 05 00 07 21 19 20 09 02
 13 05 20 01 12 12 09 03 01
 17 10 20 08 06 04 02 12 03
 21 19 00 14 15 14 00 04 09
 13 05 20 01 12 12 09 03 01
 07 24 20 15 00 26 09 07 10
 19 16 21 20 01 14 04 21 13
 13 05 20 01 12 12 09 03 01
 05 21 14 21 13 26 13 24 14

⇒ the final word is: *gjthfoblcgxtozigjemuumzmxm*

4. Let R and S be 2 relations. R has 10.000 records; a page can hold 10 R records. S has 2.000 records; a page can hold 10 S records. (1.5p)
- a. 52 buffer pages are available. Compute the cost of $R \otimes_{R.a=S.b} S$ using *page-oriented nested loops join* and *block nested loops join*; S is the outer relation.
 - b. Compute the cost of sorting R using *external merge sort* with 200 buffer pages.
 - c. R is stored at Madrid, S is stored at Athens. Compute the cost of $R \otimes_{R.a=S.b} S$ using *simple nested loops join (tuple-oriented)* in Athens, without caching; S is the outer relation.

a. $\text{SELECT } *$

$\text{FROM } R \text{ INNER JOIN } S \text{ ON } R.a = S.b$

$R - 10000 \text{ records; a page can hold } 10R \text{ records} \Rightarrow 1000 \text{ pages}$
 $S - 2000 \text{ records; a page can hold } 10S \text{ records} \Rightarrow 200 \text{ pages}$

• page oriented nested loops:

$$200 + 200 * 1000 = 200200 \text{ I/Os}$$

• block nested loops join:

block size: 50 $\Rightarrow \left\lceil \frac{200}{50} \right\rceil = 4S \text{ blocks}$

$$200 + 4 * 1000 = 4200 \text{ I/Os}$$

b. $2 * 1000 * 2 = 4000 \text{ I/Os}$

$$2 * N * \left(\left\lceil \log_{B-1} [N/B] \right\rceil + 1 \right) \text{ I/Os}$$

c. $\text{SELECT } *$

$\text{FROM } R \text{ INNER JOIN } S \text{ ON } R.a = S.b$

• t_d time to R/w a page from / to disk

• t_s time to ship a page

$$200t_d + 2000 * 1000 (t_d + t_s) = 200t_d + 2000000(t_d + t_s)$$